

# Week 6: Storage Paradigms in Cloud Computing

## NT524 — Cloud Architecture and Security

PhD. Nguyen Ngoc Tu

October 15, 2025

---

*Thuật ngữ:* storage paradigms = mô hình lưu trữ, cloud architecture = kiến trúc đám mây

# Outline

1 Multi tenant SDN review

2 Cloud storage models

- Block Storage
- Object Storage
- File Storage

3 Case study: SDN & Storage Networking

4 Hand-ons

# Learning Objectives

- Motivate storage choices for **cost optimization** and **resilience**.
- Understand **block**, **object**, **file** storage models.
- Apply **data redundancy**: replication, erasure coding, RAID.
- Design for **performance & durability** with lifecycle policies.
- Hands-on: OpenStack, AWS, Azure, GCP; backup patterns & validation.

---

*Thuật ngữ:* resilience = khả năng chịu lỗi; redundancy = dư thừa; lifecycle policy = chính sách vòng đời

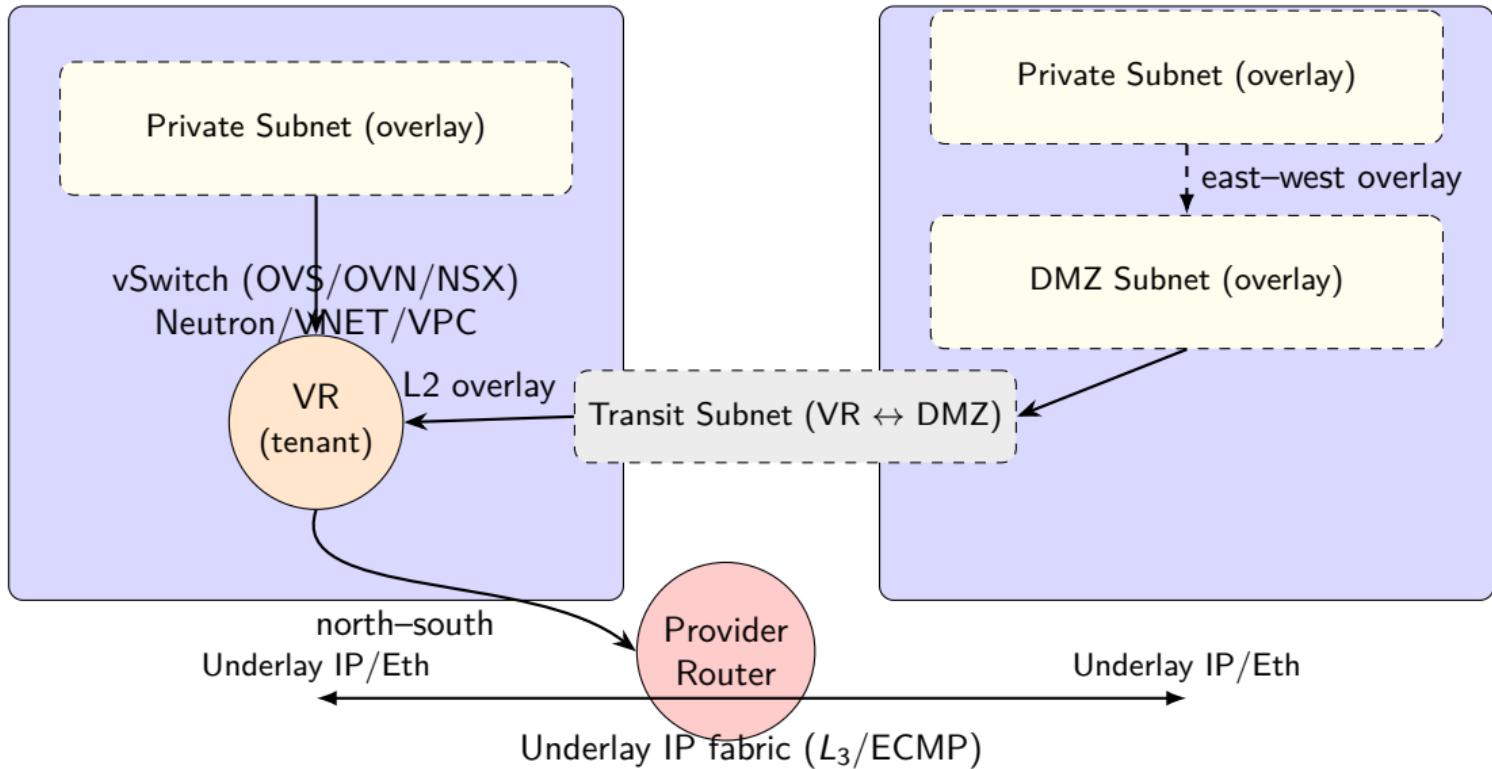
# Motivation

- **Cost:** proper tiering, compression/dedupe, lifecycle (hot/warm/cold).
- **Resilience:** RPO/RTO, multi-AZ/region, tested backups & restores.
- **Fit:** DB → block; media/backup → object; shared homes → file.

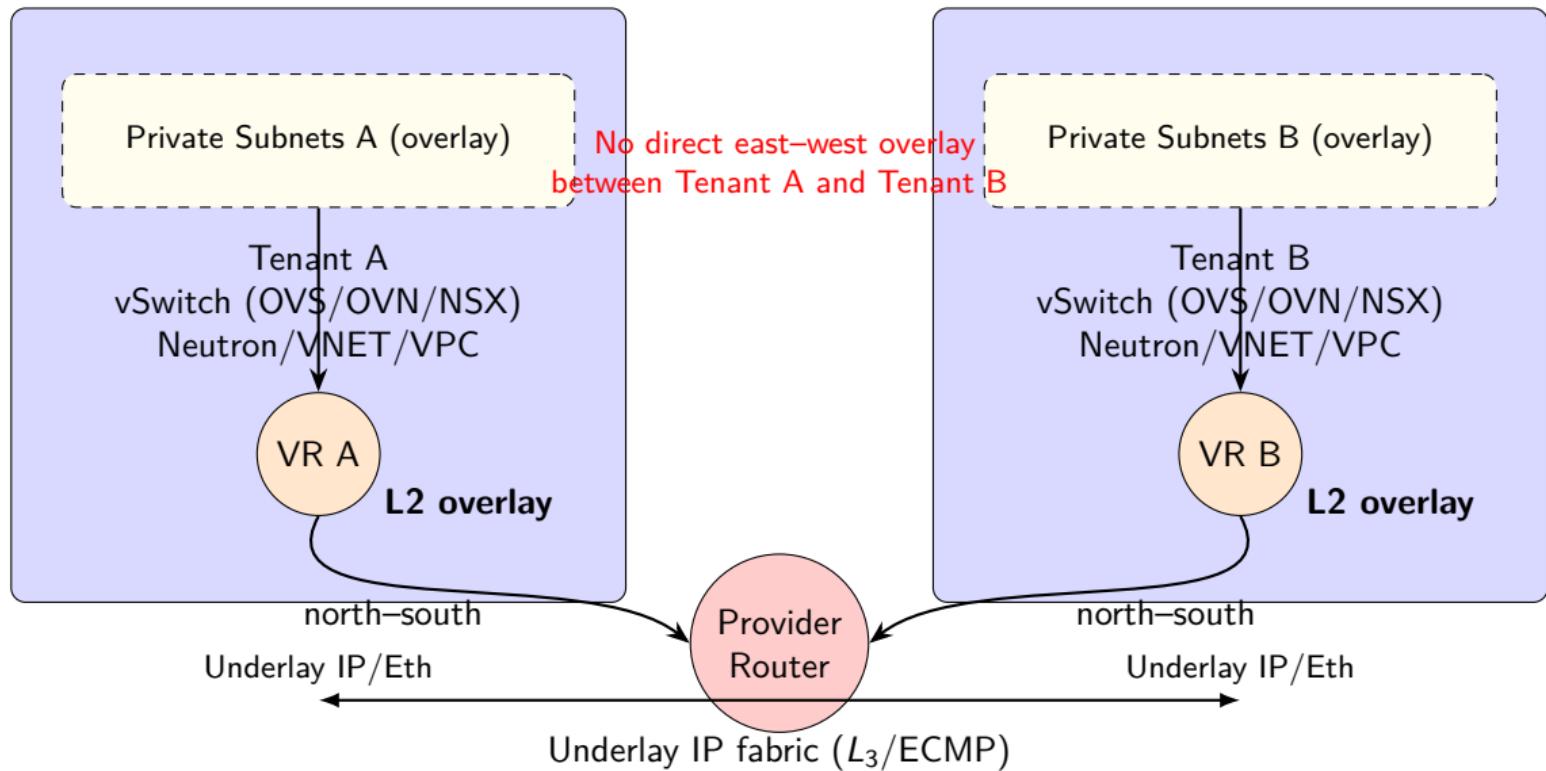
---

*Thuật ngữ: RPO/RTO = mục tiêu điểm/khôi phục; multi-AZ = đa vùng sẵn sàng*

# Big Picture: Underlay vs. Overlay (Single tenant with Transit Subnet)



# Big Picture: Underlay vs. Overlay (Multi-Tenant Isolation)



# Overlay–Underlay Mapping (Concepts)

- **Underlay:**

- Physical IP network (switches, routers, fabric).
- Provides basic **IP reachability** between nodes.
- Examples: Data center leaf-spine fabric, WAN backbone.

- **Overlay:**

- Logical network built *on top of* underlay.
- Encapsulation: **VXLAN**, **GENEVE**, NVGRE.
- Creates isolated tenant networks with their own CIDR.
- Decouples tenant addressing from physical topology.

- **Underlay ↔ Overlay Mapping:**

- Overlay packets are encapsulated inside underlay IP packets.
- **VNI (Virtual Network Identifier, 24 bits)** ↔ Underlay UDP port + IP path.
- Example: Tenant subnet 10.0.1.0/24 → VXLAN VNI 5001 → transported over underlay IP fabric.

---

Thuật ngữ: Underlay = Mạng vật lý; Overlay = mạng ảo; Mapping = ánh xạ mạng vật lý - mạng ảo

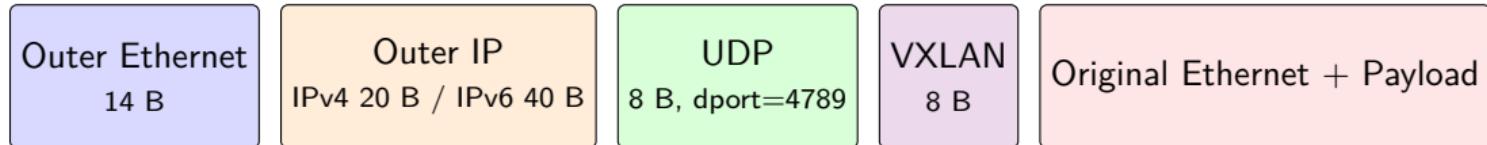
# Isolation — the Rule of Three

- ① **Separate overlay IDs:** every tenant owns a unique **VNI** and its own logical switch/router (no shared L2/L3).
- ② **Default deny across tenants:** no A↔B datapath in the vSwitch; only intra-tenant (same VNI) is reachable.
- ③ **Edge mediation if needed:** any A↔B must hairpin at the **provider edge/gateway** (FW/NAT/VRF policy).

---

*Thuật ngữ:* VNI = định danh mạng ảo; default deny = mặc định chặn; hairpin qua edge = đi vòng qua cổng biên

# VXLAN Packet Header Syntax



## Underlay (dùng để forward):

- L2: Outer Ethernet —  $dst/src$  MAC
  - L3: Outer IP —  $src/dst$  IP ( $VTEP_{src} \rightarrow VTEP_{dst}$ )
  - L4: UDP —  $src$  port (entropy),  $dport=4789$
- vSwitch/VTEP endpoints thực hiện *encapsulation* (add) và *decapsulation* (remove) các header ngoài.

## VXLAN (8 Bytes):

- 0-7: Flags ( $I=1$ )
- 8-31: Reserved
- 32-55: **VNI (24 bits)**
- 56-63: Reserved

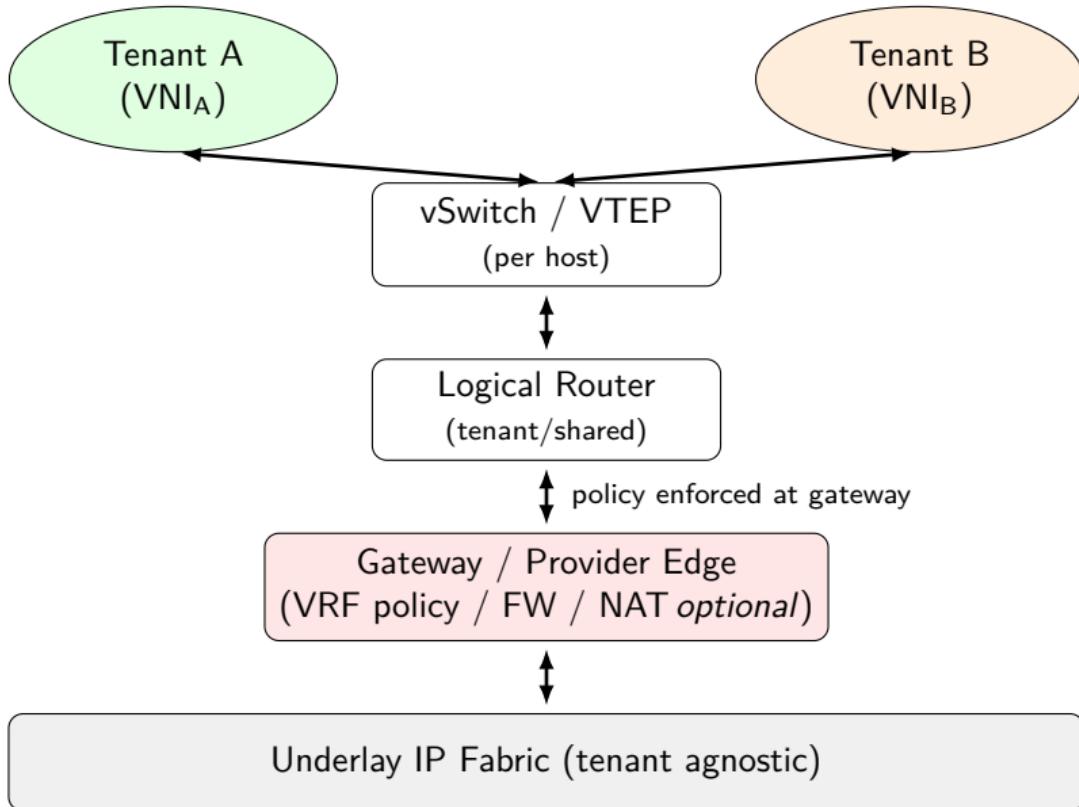
## Overlay (bảo toàn bên trong):

- Inner Ethernet —  $dst/src$  MAC
- Inner IP —  $src/dst$  IP
- Inner TCP/UDP + Payload

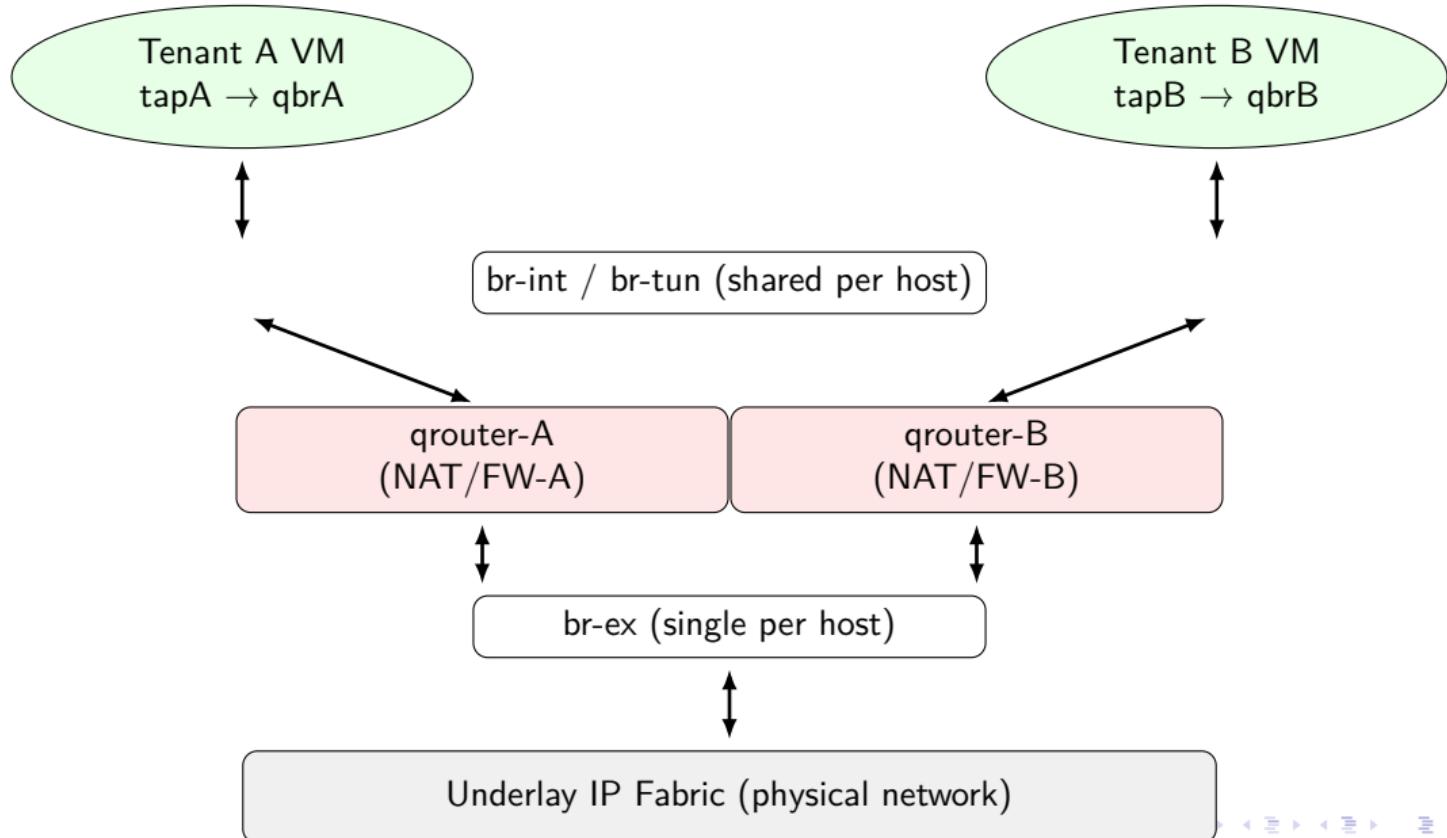
---

*Thuật ngữ:* chỉ bọc thêm header ngoài (8B, có **VNI 24-bit**);  $I=1$ . Inner (original) TCP/IP không mất.

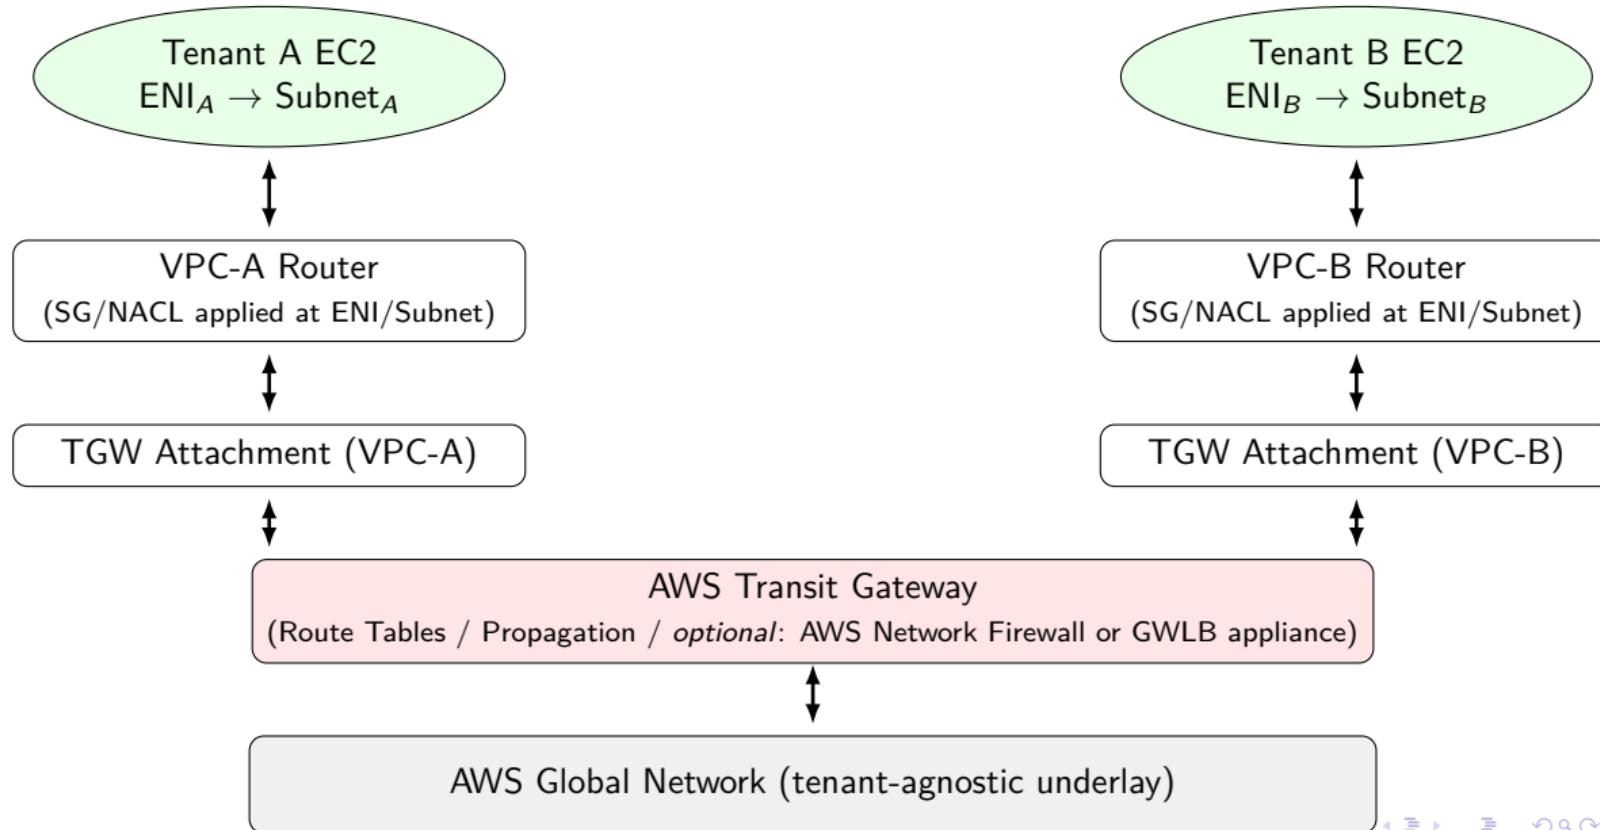
# The only allowed inter-tenant path (A $\leftrightarrow$ B via edge)



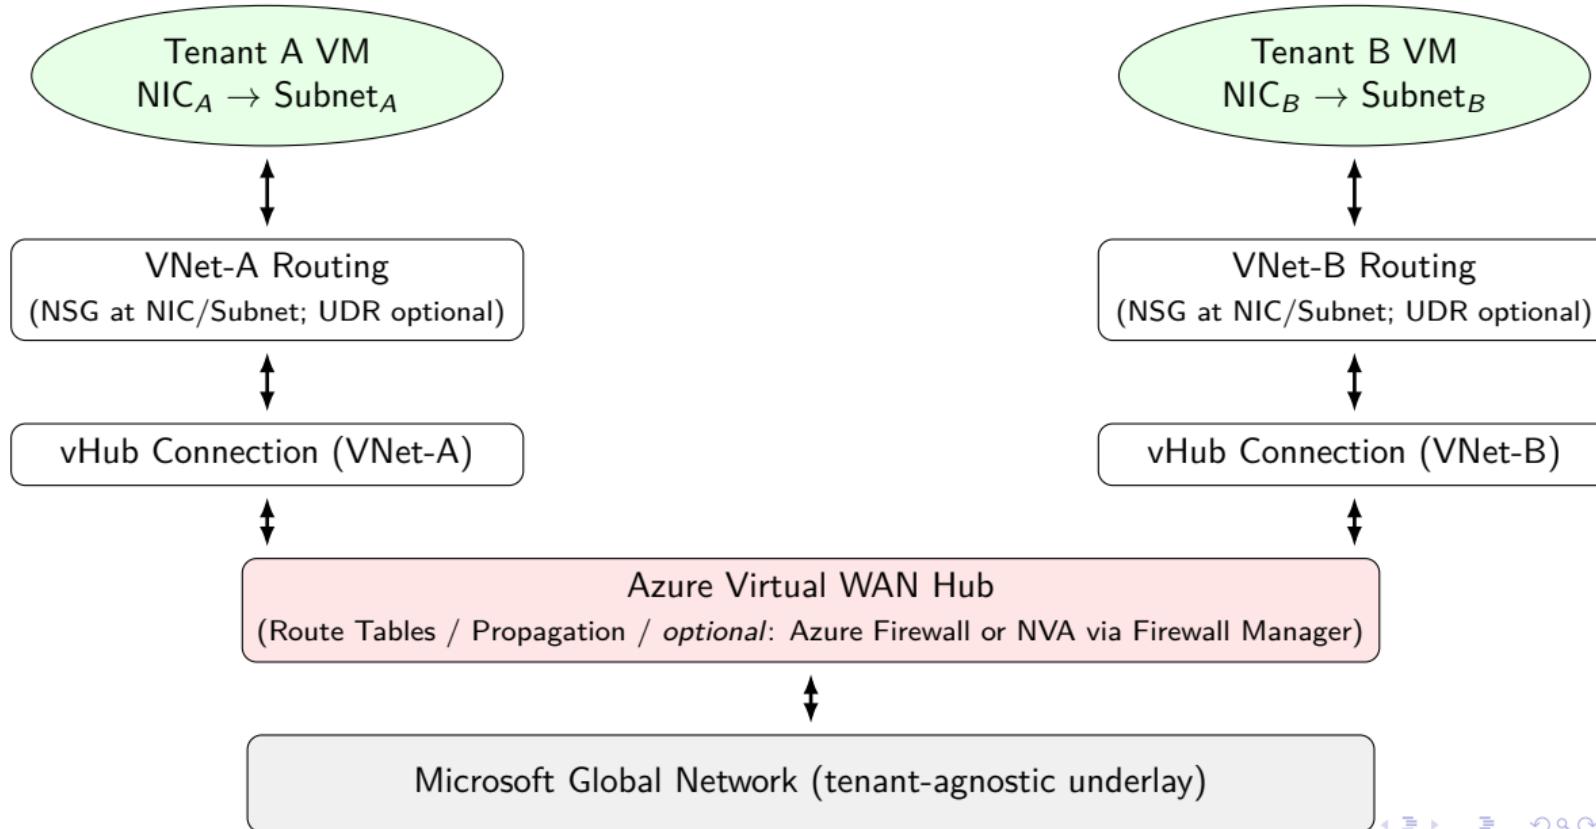
# Inter-tenant path in OpenStack (A $\leftrightarrow$ B via L3 routers)



# Inter-tenant path in AWS (A $\leftrightarrow$ B via Transit Gateway)



# Inter-tenant path in Azure (A $\leftrightarrow$ B via vWAN Hub)



# Storage Taxonomy

| Model         | Bare-metal   | Cloud cases  | Use cases   |
|---------------|--|--|---|
| <b>Block</b>  | Local SSD / NVMe namespace / SAN LUN   | VM-attached block volume (iSCSI / NVMe-oF); low latency; byte-addressable                                  | Transactional DBs, boot volumes, VM swap disks      |
| <b>Object</b> | Object array / archival system (tape library, DVD archive)-> immutability data model | Key-value object store; flat namespace; HTTP(S) API; rich metadata; <i>non-POSIX, not mounted directly</i> | Backups, media, analytics logs, data lakes          |
| <b>File</b>   | NAS filer / shared folder / network drive  | POSIX semantics via NFS/SMB; multi-client mounts; metadata ops may bottleneck                              | Home directories, content repositories, legacy apps |

Thuật ngữ: iSCSI/NVMe-oF = giao thức lưu trữ qua mạng; POSIX = chuẩn giao diện hệ thống tệp; NAS (Network-Attached Storage) = thiết bị lưu trữ gắn qua mạng.

Lưu ý: Object không mount trực tiếp; nếu dùng FUSE (s3fs/gcsfuse) thì không đầy đủ POSIX.

# Storage Access Layers: File vs Block vs Object

## • File Storage:

- Access via **file system API** (open, read, write, close).
- Provides **hierarchical directories**, file names, and POSIX semantics.
- Used for user home folders, shared drives (NFS, SMB).
- **Example:** Accessing files in a NAS share or network drive.

## • Block Storage:

- Access via **block device driver** (read/write sectors or blocks).
- OS formats the volume with a **filesystem (ext4, NTFS)** inside.
- Behaves as a raw disk; high IOPS and low latency.
- **Example:** A virtual volume attached to a VM (like /dev/vdb or C:), even if stored as a .vhdx file container.

## • Object Storage:

- Access via **HTTP(S) API** (GET, PUT, DELETE); no direct filesystem mount.
- Stores **immutable objects** = data + metadata + unique ID.
- Flat namespace (no real folders), designed for scalability and durability.
- **Example:** Uploading a file to Amazon S3 or Ceph via REST API.

---

Thuật ngữ: File: truy cập qua API của hệ tệp (thư mục, tên file); Block: truy cập qua driver thiết bị khối (sector, volume, filesystem bên trong); Object: truy cập qua API HTTP (object bất biến, không mount trực tiếp).

# Block Storage Fundamentals

- **Concept:** Data stored in fixed-size blocks (e.g., 512 B, 4 KB) with logical addressing, similar to disk sectors.
- **Model:** *Mutable data model* — blocks can be updated in place (read/write), enabling low-latency transactional workloads.
- **Architecture:** Volume service → backend (LVM, Ceph RBD [*RADOS Block Device*], NVMe-oF, iSCSI, SAN). Volumes attach to hypervisors or VMs **like virtual disks** (.vhdx, .vmdk, or .qcow2). Cloud analogues: *EBS*, *Azure Managed Disk*, *Cinder Volume*.
- **Attachment:** A block volume is typically attached to a **single host**. For **multi-attach**, a clustered filesystem (e.g., GFS2, OCFS2, VMFS on ESXi) or **read-only shared** mode is required to prevent data corruption.
- **Performance:** Defined by **IOPS** (operations per second), **throughput** (MB/s), and **latency**. Influenced by **queue depth**, **block size**, **caching policy**, and **tiering strategy**. NVMe and NVMe-oF now dominate due to microsecond-scale latency.
- **Use cases:** Databases, VM boot disks, container storage (PV in Kubernetes), snapshots and clones for rapid provisioning.

# Block Storage — Control Plane (Who Decides)

- **Lifecycle orchestration:** API → scheduler → backend driver.
  - OpenStack: cinder-api, cinder-scheduler, cinder-volume
  - Cloud parallels: EBS control plane, Azure Disk RP, GCP PD manager
- **What it handles:** provision, attach/detach, resize, migrate, snapshot policy, quotas.
- **Replication & consistency:**
  - *Synchronous* (zero RPO; higher latency) vs *asynchronous* (non-zero RPO; geo DR)
  - Quorum/journalized replication across failure domains (e.g., Ceph, ONTAP, EBS Multi-AZ)
- **Encryption & keys:**
  - Per-volume keys via KMS: Barbican (OpenStack), AWS KMS, Azure Key Vault, GCP Cloud KMS
  - Data-at-rest; attach path negotiates keys/policies
- **Governance & telemetry:** quotas, RBAC, audit; metrics/alarms via Ceilometer/CloudWatch/Azure Metrics.

---

Thuật ngữ: control plane = mặt phẳng điều khiển; RPO = mục tiêu điểm khôi phục

## Block Storage — Data Plane (Who Moves Bytes)

- **I/O path:** Guest VM → hypervisor (QEMU/KVM/ESXi) → protocol (iSCSI / NVMe-oF / RBD) → storage nodes/backends.
- **Ceph example:** librbd ↔ RADOS (striping, replication/EC).
- **Performance levers:** IOPS, throughput, latency; queue depth, block size, caching, tiering.
- **Snapshots & clones:** CoW or RoW mapping (fast, space-efficient); ideal for gold images/instant restores.
- **QoS & policy in data path:** IOPS/MBps limits, latency targets; placement/tiering hints.
- **Security in flight:** TLS/IPsec for iSCSI/NVMe-oF where required; at-rest encryption handled by backend.
- **Modern note (2025):** NVMe/NVMe-oF widely preferred for microsecond-scale latency and offload support.

---

Thuật ngữ: data plane = *mặt phẳng dữ liệu*; CoW/RoW = *ghi sao chép/ghi chuyển hướng*

# Performance Tuning & Reliability (Compact)

- **IOPS/Throughput:** tune *queue depth* (blk-mq/NVMe multi-queue), align *block size* to stripe, enable request merge for random; multi-path (iSCSI) / ANA (NVMe-oF); NUMA/IRQ pinning.
- **Latency (p95/p99):** cut underlay hops, verify PMTU; prefer NVMe-oF (RoCE/TCP) for  $\mu$ s paths; watch GC/compaction & write amplification.
- **Caching/Tiers:** host cache *read-through/write-back* (BBU/journal); SSD/NVMe hot tier + HDD cold tier; FS opts: noatime, scheduled TRIM.
- **Reliability:** spread across failure domains (disk→host→rack/AZ); RAID10/3×replica = low-latency, high space; EC( $k+m$ ) = space-efficient, higher CPU/latency, slower rebuild; sync (zero RPO) vs async (geo DR).
- **QoS/Multi-tenancy:** token-bucket min/max + **burst credits**; enforce at hypervisor/backend; fairness via cgroup `io.cost`/weights; optional DSCP for storage flows.
- **Observability/Tests:** watch queue depth, device latency, iowait, rebuild; set SLOs on p99; use `fio` with realistic `iodepth`, `numjobs`, R/W mix.

---

Thuật ngữ: PMTU = MTU hiệu dụng; ANA = Asymmetric Namespace Access; EC = mã xoá; burst credits = tín dụng bùng nổ.

# Provider Implementations (Well-known Cloud Examples)

| Provider            | Tiers / Types   | Key features & notes   |
|---------------------|---|--|
| AWS EBS             | gp3, io1/io2, <i>io2 Block Express</i>                              | AZ-scoped volumes; provisioned IOP-S/throughput (gp3 decouples); Multi-Attach (io1/io2 on Nitro); snapshots to S3; encryption via KMS. |
| Azure Disks         | Standard HDD/SSD, Premium SSD, <i>Premium SSD v2</i> , Ultra Disk   | Managed disks; zonal/regional options; Disk Encryption Sets; shared disks for clustering; bursting on select SKUs.                     |
| GCP Persistent Disk | Standard, Balanced, SSD; <i>Hyper-disk (Extreme/Throughput)</i>     | Live resize; regional PD; multi-writer (RWX) for select types; encryption by default / CMEK.   |
| OpenStack Cinder    | LVM, Ceph RBD, NetApp ON-TAP, Dell EMC (PowerFlex/PowerStore), etc. | API/scheduler/driver control plane; multi-attach since Train; Barbican encryption; per-backend QoS/policy.                             |
| VMware vSAN         | Clustered NVMe/SSD/HDD pool (policy-based)                          | Storage policies (RAID 1/5/6, FTT); dedupe/-compression; stretched cluster; vSphere integration.                                       |

# Emerging Trends & Integration with Cloud Services

- **Disaggregated Storage:** NVMe-oF / RoCE / EBOF (Elastic Block Over Fabric) decouples compute.
- **Smart NIC offload:** DPUs handle I/O and encryption, reducing host latency.
- **Container Integration:** CSI drivers (Cloud Storage Interface) for Kubernetes.
- **Hybrid Backups:** Snapshots → Object Storage (S3, Azure Blob, GCS).
- **Next Wave:** zonal replica pools, I/O telemetry for AIOps, ML-based tier optimization.

---

Thuật ngữ: disaggregated storage = lưu trữ tách rời tính toán; DPU = bộ xử lý dữ liệu mạng

# Object Storage Basics

- **Model:** bucket/container → object (key, data, metadata).
- **Scale:** scale-out via sharding/partitioning; billions of objects; large data (logs, media, backups).
- **Features:** pre-signed URL, lifecycle tiers, versioning, server-side encryption.
- **Access:** HTTP/HTTPS (REST); strong API contracts; eventual &/or strong read-after-write consistency (provider-specific).

---

*Thuật ngữ:* pre-signed URL = URL đã ký; server-side encryption = mã hóa phía máy chủ

# Internal Architecture & Consistency

- **Frontends:** stateless API gateways/routers; request auth (sigv4/SAS/OAuth).
- **Metadata plane:** namespace index, bucket policies, version maps, object manifests.
- **Data plane:** erasure coding or replicated chunks across failure domains (AZs).
- **Placement:** partition by key prefix/hash; hot partitions auto-split/merge.
- **Consistency:** modern clouds offer **strong** read-after-write for new puts/overwrites; cross-region replication may be async.

---

*Thuật ngữ:* erasure coding = mã hóa xoá; failure domain = miền lỗi

# Performance Patterns & Cost Controls

- **Throughput:** parallelism + multi-part upload; range GET for partial reads.
- **Key design:** random/hashed prefixes to avoid hot partitions; avoid sequential keys.
- **Latency classes:** frequent access vs infrequent vs archive; retrieval SLA differs.
- **Cost levers:** storage class, lifecycle policies, intelligent tiering/auto-class, compression.
- **Data access accel:** edge caches/CDN; S3 Select/Blob Query for in-place filtering.

---

*Thuật ngữ:* multi-part upload = tải nhiều phần; hot partition = phân vùng nóng

# Provider Implementations (Well-known Examples)

- **AWS S3:** Standard/IA/One Zone-IA/Glacier (Instant/Flexible/Deep Archive), Intelligent-Tiering; Access Points; CRR/SRR.
- **Azure Blob:** Hot/Cool/Archive; immutable policies (time-based/legal hold); SAS; hierarchical namespace (Data Lake Gen2).
- **GCP Cloud Storage:** Standard/Nearline/Coldline/Archive; Autoclass; Turbo/dual-region options; uniform IAM.
- **OpenStack Swift:** proxy + account/container/object rings; EC or replication; TempURL; middleware pipeline.
- **S3-compatible (on-prem):** MinIO, Ceph RGW; CSI for Kubernetes; tiering to public clouds.

---

*Thuật ngữ: SAS = chữ ký truy cập chia sẻ; CRR/SRR = sao chép chéo vùng/cùng vùng*

# Security, Governance & Emerging Trends

- **Identity/Policy:** IAM + bucket/container policies; block public access; VPC endpoints/PrivateLink.
- **Encryption:** SSE with provider KMS or customer-managed keys; client-side encryption for end-to-end.
- **Data protection:** Versioning, Object Lock (WORM), replication (same/cross-region), integrity checks (ETag/MD5/SHA-256).
- **Eventing:** native notifications to queues/bus (EventBridge/Event Grid/Pub/Sub) for data pipelines.
- **Trends:** storage-lake convergence (lakehouse), multi-region active/active, governance catalogs, zero-ETL analytics on objects.

---

*Thuật ngữ: WORM = ghi một lần, đọc nhiều lần; PrivateLink = kết nối riêng qua mạng nhà cung cấp*

# File Storage Concepts

- **Protocols:** NFSv3/v4.1/v4.2, SMB 2/3; POSIX semantics, advisory/mandatory locking.
- **Namespace:** hierarchical paths, directories, ACLs, quotas.
- **Performance:** throughput vs concurrency; metadata ops (stat/readdir) are bottlenecks.
- **Use cases:** home dirs, web content, render farms, EDA/HPC, legacy apps.
- **Access modes:** multi-client shared (RWX), file/byte-range locking, close-to-open consistency (NFS).

---

*Thuật ngữ:* locking = khoá tệp; metadata ops = tác vụ siêu dữ liệu; RWX = đọc-ghi-thực thi dùng chung

# Internal Architecture & Scale-Out NAS

- **Control plane:** export/share policies, authN/Z (LDAP/AD, Kerberos), snapshots.
- **Data plane:** NAS heads → backends (RAID/erasure) via SAS/NVMe; scale-out adds *multiple* heads.
- **Metadata:** single vs distributed MDS; directory sharding; journaling.
- **Parallelization:** pNFS layouts (files/objects/blocks); SMB Multichannel; NFS over RDMA.
- **Multi-protocol:** NFS+SMB to same dataset (name-mapping, ACL translation).

---

Thuật ngữ: MDS = máy chủ siêu dữ liệu; pNFS = NFS song song; RDMA = truy cập bộ nhớ từ xa

# Performance, Tuning & Operations

- **Throughput vs IOPS:** large sequential IO vs small random; many small files hurt.
- **Client tuning:** NFS rsize/wsize, nconnect, attribute caching; SMB Multichannel, credits.
- **Latency:** network RTT, server CPU, metadata locks; prefer close clients/AZ locality.
- **Hotspots:** single directory hot keys; hash/fan-out directory design.
- **Ops:** snapshots/clones, quotas (user/group/project), tiering to object, backups.

---

Thuật ngữ: snapshot = bản chụp nhanh; fan-out = toả nhánh

# Provider Implementations (Well-known Examples)

- **AWS EFS:** regional, NFSv4.1; classes: Standard/IA; throughput modes: Bursting/Provisioned; EFS-to-S3 lifecycle.
- **AWS FSx:** NetApp ONTAP (NFS/SMB/SnapMirror), Lustre (HPC), OpenZFS; multi-AZ options.
- **Azure Files:** SMB/NFS; tiers Hot/Cool/Premium; **ANF** (Azure NetApp Files) for high IOPS/low latency.
- **GCP Filestore:** Basic/Enterprise/High Scale; NFSv3/v4.1; regional HA in Enterprise.
- **OpenStack Manila:** drivers for CephFS, NetApp, GlusterFS, Dell/EMC, etc.; share networks, access rules.
- **HPC/On-prem:** Lustre, IBM Spectrum Scale (GPFS), BeeGFS, Isilon/PowerScale, NetApp ONTAP.

---

Thuật ngữ: HA = tính sẵn sàng cao

# Security, Governance & Advanced Patterns

- **Identity/ACLs:** POSIX vs NFSv4/SMB ACLs; AD/Kerberos; least privilege shares/exports.
- **Encryption:** at-rest (KMS/Key Vault/Cloud KMS) & in-flight (TLS, SMB encryption, Kerberos signing).
- **Data protection:** snapshots, replication (sync/async, cross-AZ/region), ransomware-resilient immutable snaps.
- **Kubernetes:** CSI drivers for RWX; dynamic provisioning; workload-affinity to reduce latency.
- **Trends:** disaggregated NAS over NVMe-oF, DPU offload for crypto/IO, multi-protocol data lakes, policy-driven tiering to object.

---

*Thuật ngữ:* immutable snapshot = bản chụp không thể sửa; CSI = giao diện lưu trữ cho container

## Case Setup & SLOs (Cross-Cloud)

- **Workloads & I/O:** DB (8–64 KB random RW), media render (1–8 MB seq), analytics (mixed; large scans).
- **Protocols:** iSCSI / NVMe-oF (block), NFSv4.1/4.2 & SMB 3.x (file).
- **SDN overlay:** VXLAN/Geneve overhead ( $\approx$ 50–60 B)  $\Rightarrow$  plan MTU/headroom; rely on PMTU discovery.
- **Routing fabric:** east–west via L3 gateways (qrouter/TGW/vWAN/Cloud Router) with policy; avoid L2 extension.
- **SLO targets:** p99 latency (block:  $\leq$ 2–5 ms; file:  $\leq$ 5–15 ms), sustained throughput (GB/s), durability (nines), RPO/RTO.
- **Policy/QoS:** DSCP marking at host; per-class queues; IOPS/MBps limits where supported.

---

Thuật ngữ: PMTU = phát hiện MTU đường đi; east–west = luồng nội bộ giữa VNET/VPC/VN

# OpenStack (OVN) — iSCSI/NFS/NVMe-oF

- **Overlay/MTU:** ML2/OVN (Geneve). Set tenant MTU (e.g., 1450/9000-overhead) and underlay MTU consistently; verify with `tracepath` / `ping -M do`.
- **Cinder (block):** iSCSI or Ceph RBD. For iSCSI: `dm-multipath`, `no_path_retry`, `udev` rules; for NVMe-oF: enable NVMe multipath.
- **Manila (file):** NFS/SMB exports via share networks; align export subnet with router namespaces; security groups → minimal, allow only storage flows.
- **Routing/policy:** Inter-tenant via logical routers → gateway chassis (NAT optional). OVN ACLs at LR/LSP; ECMP for scale.
- **QoS/DSCP:** Neutron QoS (rate limit) + OVN QoS rules; DSCP mark on instances (`tc/iptables mangle`) → preserve in underlay if supported.
- **Validate:** `fio` (iodepth sweep), `iperf3` (parallel flows), `ovn-trace` (policy path), `ovn-appctl` latency.

---

Thuật ngữ: ECMP = cân bằng tải đa đường; LR/LSP = router/port logic

# AWS (Transit Gateway) — EFS/FSx & Hybrid iSCSI

- **File services:** EFS (NFSv4.1) for shared POSIX; FSx for ONTAP (NFS/SMB/iSCSI features) or FSx for Lustre (HPC).
- **Block paths:** EBS is local to EC2 (no network). For SAN-like iSCSI, use FSx/ONTAP or on-prem SAN via DX/VPN.
- **SDN path:** VPC ↔ TGW (route tables per attachment); insert Network Firewall/GWLB for policy.
- **MTU/PMTU:** Use provider defaults per ENI/VPC; verify with tracepath across TGW/ENI and DX/VPN to avoid fragmentation.
- **QoS/DSCP:** Mark at EC2; preserve across appliances; use policing on NVA if needed (not native in VPC fabric).
- **Tuning:** EFS throughput mode (Bursting/Provisioned), concurrency via multiple NFS clients; FSx/ONTAP: junctions, aggregates, FlexVol/QoS.
- **Validate:** fio on EFS mount (nconnect), nfsstat, CW metrics (PercentIOLimit), mtr across TGW.

---

Thuật ngữ: DX = Direct Connect; GWLB = cân bằng tải công

# Azure (vWAN/Hub-VNet) — ANF & Azure Files

- **File services:** **Azure NetApp Files (ANF)** for high IOPS/low latency (NFS/SMB); **Azure Files** (SMB/NFS) for broad compatibility.
- **SDN path:** Spoke VNets  $\leftrightarrow$  **vWAN Hub** or Hub-VNet; centralized policy via **Azure Firewall/NVA**; UDR for steering.
- **Identity/ACL:** AD/Kerberos for SMB; POSIX modebits/ACLs for NFS; carefully map IDs across domains.
- **MTU/PMTU:** Keep consistent per NIC/vNIC and gateways; validate PMTU on ER/VPN paths; avoid L2 stretch, prefer routed.
- **QoS:** ANF capacity pools/volume QoS; Azure Files tiers (Hot/Cool/Premium); traffic shaping on NVA.
- **Validate:** fio over NFS/SMB, nfsstat/smbstatus, Connection Monitor, psping for tail latency.

---

Thuật ngữ: UDR = định tuyến do người dùng; ER = ExpressRoute

# GCP (Cloud Router) — Filestore & Hybrid NVMe-oF

- **File services:** **Filestore** (NFSv3/v4.1) tiers: Basic/Enterprise/High Scale; regional HA in Enterprise.
- **Block/high-perf:** For extreme IOPS, consider local SSD or Hyperdisk (for DB) + app-level sharding; NVMe-oF usually via partner NVAs.
- **SDN path:** VPCs via **VPC peering** or **Cloud Router** across **Cloud Interconnect/VPN**; centralize policy with firewalls/NVAs.
- **MTU/PMTU:** Keep per-VPC/NIC consistent; verify PMTU across Interconnect; avoid fragmentation on storage flows.
- **QoS/DSCP:** Mark at VM; enforce in NVAs; use Filestore performance guidance (num threads, rsize/wsize).
- **Validate:** fio profiles (randread/randwrite, bs=4k/64k/1m), iperf3 with multiple streams, traceroute/tracepath.

---

*Thuật ngữ:* Interconnect = kết nối chuyên dụng tới GCP; NVA = thiết bị ảo mạng

# Data Redundancy

- **Replication (x3):** improves availability and reads; higher cost.
- **Erasure coding ( $k+m$ ):** capacity efficiency =  $k/(k + m)$ ; overhead =  $(k + m)/k$ ; best for cold/large objects.
- **RAID:** 0/1/5/6/10; parity write penalties; rebuild risk windows.
- **Durability vs availability:** data loss vs uptime.

---

*Thuật ngữ:* erasure coding = mã xoá; availability = khả dụng

# Cloud Mapping (OpenStack / AWS / Azure / GCP)

| Need   | OpenStack              | AWS                              | Azure / GCP                        |
|--------|------------------------|----------------------------------|------------------------------------|
| Block  | Cinder (LVM/Ceph)      | EBS (gp3, io1/io2; Multi-Attach) | Managed Disks / PD (balanced/ssd)  |
| Object | Swift (replication/EC) | S3 (strong consistency)          | Blob / GCS (Std/Near-Cold/Archive) |
| File   | Manila (NFS/SMB)       | EFS / FSx                        | Azure Files (SMB/NFS) / Filestore  |
| Keys   | Barbican               | KMS                              | Key Vault / KMS                    |
| Backup | Snapshots → object     | EBS snapshots → S3               | Snapshots → Blob/GCS               |

Thuật ngữ: strong consistency = nhất quán mạnh; Filestore = dịch vụ NFS GCP

# Hands-on: OpenStack

**Goal:** Cinder/Swift/Manila via Ansible; snapshot & verified restore.

- ① Create *volume type* + *QoS*; attach to VM; format/mount; run **fio**.
- ② Swift container; upload; set X-Delete-After; verify expiry.
- ③ Manila NFS share; IP access; mount on 2 VMs; test locks/concurrent writes.
- ④ Snapshot volume; archive to Swift; **restore** to new VM; verify SHA-256.

---

*Thuật ngữ:* QoS = chất lượng dịch vụ; expiry = hết hạn; SHA-256 = băm kiểm toàn vẹn

# Ansible: Volume Type + QoS (OpenStack)

```
- hosts: controller
  tasks:
    - name: Create QoS and volume type, then associate
      openstack.cloud.qos:
        state: present
        name: ssd-qos
        specs:
          total_iops_sec: "3000"

    - name: Create volume type
      openstack.cloud.volume_type:
        state: present
        name: ssd-qos-type

    - name: Associate QoS to type
      openstack.cloud.volume_type:
        state: present
        name: ssd-qos-type
        qos_specs: ssd-qos

    - name: Create 20GB volume using the type
```

# Manila NFS: Export & Mount (OpenStack)

```
- hosts: controller
  tasks:
    - name: Create NFS share
      openstack.cloud.share:
        state: present
        name: wk6-share
        size: 10
        share_proto: NFS
        is_public: false
      register: create_share

    - name: Get export location
      openstack.cloud.share_info:
        name: wk6-share
      register: share_info

    - set_fact:
        wk6_share_export: "{{ share_info.shares[0].export_locations[0].path }}"

- hosts: apps
  become: true
```

# Swift Lifecycle & Backup (OpenStack)

```
- hosts: controller
  tasks:
    - name: Create container with metadata
      openstack.cloud.object_container:
        state: present
        name: wk6-backups
        metadata: { "retention": "180d" }

    - name: Upload with expiry (X-Delete-After)
      openstack.cloud.object:
        state: present
        container: wk6-backups
        name: snap-{{ ansible_date_time.epoch }}.raw
        filename: /tmp/snap.raw
        headers:
          X-Delete-After: "15552000" # 180 days

    - name: Create snapshot
      openstack.cloud.volume_snapshot:
        state: present
        name: wk6-db-snap-{{ ansible_date_time.date }}_11
```

# Hands-on: AWS

**Goal:** EBS (gp3/io2, Multi-Attach), S3 (strong consistency), EFS.

- ① Create EBS gp3; attach to EC2; benchmark fio; try Multi-Attach with a cluster FS.
- ② S3 bucket: enable versioning; pre-signed URL; lifecycle to Glacier.
- ③ EFS: mount from 2 EC2; test locks; observe CloudWatch metrics.

---

*Thuật ngữ:* Multi-Attach = gắn nhiều phiên; Glacier = lưu trữ lạnh

# Hands-on: Azure

**Goal:** Managed Disks, Blob, Azure Files (SMB/NFS, Premium).

- ① Managed Disk: attach to VM; `fio gp/premium`; monitor IOPS.
- ② Blob: container, SAS tokens, lifecycle (cool/archive); test restore.
- ③ Azure Files: SMB or NFS (Premium); try SMB Multichannel if supported.

---

*Thuật ngữ:* SAS = *chữ ký chia sẻ*; Multichannel = *đa kênh SMB*

# Hands-on: GCP

**Goal:** Persistent Disk (pd-balanced/pd-ssd), Cloud Storage (Std/Near/Cold/Archive), Filestore.

- ① PD: attach to GCE; run `fio`; snapshots.
- ② Cloud Storage: classes; signed URL; lifecycle; check retrieval cost.
- ③ Filestore: NFS mount from 2 VMs; test POSIX locks.

---

*Thuật ngữ:* GCE = máy ảo GCP; PD = đĩa bền

# Performance & Cost

- **Benchmark:** fio 4k randread (QD=1/16/32), 4k 70/30, 1M sequential; record p50/p95/p99.
- **Tiers:** hot/warm/cold; compression; columnar formats (Parquet/ORC) for analytics.
- **Egress/requests:** batch, CDN, concurrency.
- **FinOps:** track /TB-month, /IOPS; prune orphaned snapshots.

---

Thuật ngữ: p50/p95/p99 = bách phân vị; egress = dữ liệu ra

# Security & Compliance

- **At rest:** LUKS, backend-native, KMS/Barbican; per-project keys.
- **In transit:** TLS/mTLS; NFS/SMB encryption where supported; IPsec for iSCSI if needed.
- **Access:** RBAC/ABAC; bucket policies; block public access; least privilege.
- **Immutability:** WORM/object lock; protect snapshots/backups from deletion.

---

Thuật ngữ: ABAC = kiểm soát theo thuộc tính; WORM = ghi một lần đọc nhiều lần

# Threat Model (STRIDE) for Storage

- **Spoofing:** unauthorized bucket access ⇒ IAM/RBAC, short-lived pre-signed URLs.
- **Tampering:** object/volume modification ⇒ versioning, WORM, checksum, signatures.
- **Repudiation:** missing audit ⇒ enable logs (Cinder/Manila/Swift/EBS/S3/EFS).
- **Information Disclosure:** public buckets, exposed snapshots ⇒ policies & encryption.
- **DoS:** IO storms ⇒ QoS limits; API rate limits.
- **EoP:** container/hypervisor escape ⇒ patching, network policy, isolated IAM.

---

Thuật ngữ: audit = ghi nhật ký kiểm toán; rate limit = giới hạn tốc độ yêu cầu

# Decision Flow

- ① **Access pattern?** random vs sequential; small vs large objects.
- ② **Consistency?** strong vs eventual; single vs multi-region.
- ③ **Failure domain?** disk/node/AZ/region.
- ④ **SLO & Cost?** p95 latency, /TB-month, /IOPS.
- ⑤ **Governance?** PII, retention, immutability, legal hold.

---

Thuật ngữ: failure domain = miền lỗi; governance = quản trị dữ liệu

# Comparison: Services & Features

|             | <b>OpenStack</b>        | <b>AWS</b>                   | <b>Azure</b>                | <b>GCP</b>                  |
|-------------|-------------------------|------------------------------|-----------------------------|-----------------------------|
| Block       | Cinder (QoS, snapshots) | EBS (gp3, io2, Multi-Attach) | Managed Disks (Std/Premium) | PD (standard/balanced/ssd)  |
| Object      | Swift (replication/EC)  | S3 (strong R/W/LIST)         | Blob (hot/cool/archive)     | GCS (Std/Near-Cold/Archive) |
| File        | Manila (NFS/SMB)        | EFS/FSx                      | Azure Files (SMB/NFS)       | Filestore (NFS)             |
| Consistency | Backend-dependent       | Strong                       | Strong (Blob)               | Strong reads (ops vary)     |
| Keys        | Barbican                | KMS                          | Key Vault                   | KMS                         |
| Notes       | OVN SDN common          | Nitro + NVMe                 | SMB Multichannel            | Regional/multi-reg buckets  |

*Thuật ngữ:* LIST = liệt kê; Nitro = ảo hóa phần cứng AWS; Multichannel = đa kênh SMB

# Assessment & Deliverables

- **Report (2–3 pages):** architecture for 3 workloads; technical & cost rationale; clear SLOs.
- **Evidence:** screenshots: volumes/attach, buckets/containers, share mounts, fio, restores.
- **Files:** playbooks, scripts, lifecycle, recovery runbook.
- **Rubric (10 pts):** functionality (4), performance (2), security (2), docs (2).

---

*Thuật ngữ: runbook = sổ tay quy trình; rubric = thang điểm chi tiết*

# Wrap-up

- Match block/object/file to workload and budget.
- Choose redundancy wisely (replication/EC/RAID), encrypt, apply lifecycle.
- SDN impacts storage data paths: MTU, QoS, multipath.
- Next: Container Fundamentals.

---

*Thuật ngữ:* data plane = *mặt phẳng dữ liệu*; control plane = *mặt phẳng điều khiển*