

Week 7: Container Fundamentals

NT524 — Cloud Architecture and Security

PhD. Nguyen Ngoc Tu

October 22, 2025

Thuật ngữ: container fundamentals = nền tảng container; cloud architecture = kiến trúc đám mây

Outline

- 1 Review on SDN and storage models
- 2 Cloud Application Isolation and Deployment
- 3 Container Runtime Engines
- 4 Container Platform

Learning Objectives

- Application Isolation models
- Containers and VMs: Compare and contrast
- Container usage: portability, lower overhead, microservices
- Container building blocks: container **runtime** and **platform** (*orchestrator boundary only; details in Week 8*)

Thuật ngữ: platform = nền tảng container; orchestrator = bộ điều phối (đến tuần sau)

Review — Isolation Mindset (Compute → Network → Storage)

- **Compute / Hardware (VMs):** Hypervisor isolation; tenancy boundaries; vCPU/NUMA pinning; I/O isolation (IOMMU, SR-IOV) to contain blast radius and noisy neighbors.
- **Network (SDN):** Tenant segmentation (VRF/VLAN/VXLAN/Geneve), Security Groups/ACLs, microsegmentation, controlled exposure (LB/NAT), and policy enforcement.
- **Storage:** Ephemeral vs persistent; block/object/shared filesystems; snapshots/clones; encryption-at-rest; access modes (RWO/RWX); quotas & I/O QoS to bound interference.

Thuật ngữ: VRF (Virtual Routing and Forwarding) = *định tuyến & chuyển tiếp ảo*; VLAN (Virtual LAN) = *mạng LAN ảo*; VXLAN/Geneve = *mạng phủ đóng gói L2 trên L3*; Security Group = *nhóm bảo mật*; ACL (Access Control List) = *danh sách kiểm soát truy cập*; microsegmentation = *vi phân đoạn*; IOMMU (Input–Output Memory Management Unit) = *bộ ánh xạ bộ nhớ I/O*; SR-IOV (Single Root I/O Virtualization) = *ảo hoá I/O đơn gốc*; LB/NAT (Load Balancer / Network Address Translation) = *cân bằng tải / biên dịch địa chỉ mạng*; RWO/RWX (ReadWriteOnce / ReadWriteMany) = *đọc–ghi độc quyền / đọc–ghi chia sẻ*; QoS (Quality of Service) = *chất lượng dịch vụ*

Motivation — Application Isolation (Why we care)

- **Contain failures and cross-tenant impact:** isolate each service/tenant so one compromise or crash cannot cascade.
- **Portable, safe artifacts:** ship one minimal, reproducible, signed image that behaves the same across environments.
- **Enable multi-tenant PaaS on one host:** enforce resource limits and admission policy to prevent interference and keep SLOs predictable.

Thuật ngữ: multi-tenant = *đa thuê*; non-root = *không quyền root*; LSM = *mô-đun bảo mật Linux*; SLO (Service Level Objective) = *mục tiêu mức dịch vụ*

Review: Software-Defined Networking (SDN) in Cloud Platforms

- **Concept:** SDN separates the **control plane** (policy/intent) from the **data plane** (packet forwarding). In cloud, this is implemented by **Neutron + OVN/OVS**.
- **Control Plane:** define tenant networks, subnets, routers, ACL/QoS via APIs; store intent in OVN NB DB.
- **Data Plane:** virtual switches (OVS) and tunnels (VLAN/VXLAN/GRE/Geneve) stitch compute nodes.
- **Why it matters:** isolation between tenants/projects; programmable connectivity for any workload model.

Thuật ngữ: control plane = mặt phẳng điều khiển; data plane = mặt phẳng dữ liệu; OVN/OVS = Open Virtual Network / Open vSwitch; ACL/QoS = kiểm soát truy cập/chất lượng dịch vụ

Review: Cloud Storage Models (Block / Object / File)

Model	Access & Protocol	Typical Services	Strengths	Limitations
Block	Attached volume to host; iSCSI/NVMe	Cinder, EBS, PD, Azure MD	Low latency; consistent IOPS; boot/system disks, DBs	Per-host attach; AZ affinity; snapshot lifecycle mgmt
Object	HTTP(S) key-value objects; multipart	Swift, S3, GCS, Azure Blob	Massive scale; durability (e.g. “11 nines”); lifecycle tiers	Higher latency; eventual consistency (class-dependent); non-POSIX
File	Shared NFS/SMB namespace; multi-client mount	Manila, EFS/FSx, Azure Files, File-store	Shared tree; lift-and-shift; simple app portability	Throughput/latency tier-bound; scaling limits; cost at scale

Thuật ngữ: block/object/file = khối/đối tượng/tệp; IOPS = số thao tác I/O/giây; AZ = khu vực sẵn sàng; POSIX = chuẩn giao diện hệ tệp

Review OpenStack case study: Block / Object / File Storage

- **Block (Cinder)**: low-latency *volumes* for OS/DB; snapshots for backup/clone.
- **Object (Swift)**: HTTP objects for backups/logs/media/data lake; versioning + lifecycle.
- **File (Manila)**: shared NFS/SMB for multi-host app state or legacy lift-and-shift.
- **Operator model**: types/classes, quotas, retention, replication/QoS policies.

Thuật ngữ: retention = lưu giữ; replication = sao chép; quota = hạn ngạch

Review OpenStack case study: Cinder (Block Storage)

Goal: create, attach, and mount a persistent block volume to a VM.

1) Inspect volume types & quotas

```
openstack volume type list  
openstack limits show --absolute | grep VOLUME
```

2) Create a 10GiB volume (choose a type)

```
openstack volume create --size 10 --type lvmdriver --availability-zone nova demo-vol  
openstack volume show demo-vol -f table
```

3) Attach to a server and format/mount in-guest

```
openstack server add volume <VM_NAME> demo-vol  
# Inside VM:  
lsblk  
sudo mkfs.ext4 /dev/vdb  
sudo mkdir -p /data && sudo mount /dev/vdb /data
```

4) Persist mount (optional)

```
echo "/dev/vdb /data ext4 defaults 0 0" | sudo tee -a /etc/fstab
```

Review OpenStack Swift — Public URL (readable container)

Goal: create a container, upload/download objects, and fetch via a public URL.

1) Create a container and list

```
openstack container create demo-container  
openstack container list
```

2) Upload & list objects

```
openstack object create demo-container report.pdf  
openstack object list demo-container
```

3) Make container publicly readable (demo use)

```
openstack container set --public demo-container
```

4) Construct & fetch the URL

```
SWIFT_URL=$(swift stat -v | awk -F': ' '/StorageURL/ {print $2}')  
curl -I "$SWIFT_URL/demo-container/report.pdf"
```

Thuật ngữ: public read=công khai đọc; endpoint=điểm cuối

Review OpenStack Swift — Temp URL (time-limited)

Goal: Keep container private but share a signed, time-limited URL.

1) Set an account-level Temp-URL key

```
KEY=$(openssl rand -hex 16)
openstack object store account set --property Temp-URL-Key="$KEY"
```

2) Gather account storage URL

```
ACCOUNT=$(swift stat -v | awk -F': ' '/Account/ {print $2}')
SWIFT_URL=$(swift stat -v | awk -F': ' '/StorageURL/ {print $2})
```

3) Generate a signed URL (valid 3600s)

```
SIG=$(swift tempurl GET 3600 "/v1/$ACCOUNT/demo-container/report.pdf" "$KEY")
```

4) Fetch via Temp URL

```
curl -I "$SWIFT_URL/demo-container/report.pdf?$SIG"
```

Thuật ngữ: Temp URL=liên kết tạm thời; signed URL=liên kết có chữ ký

Review OpenStack case study: Manila (File Storage)

Goal: create an NFS share, allow access, mount from a VM.

1) Inspect share types

```
openstack share type list
```

2) Create a 10GiB NFS share

```
openstack share create nfs 10 --name demo-share  
OpenStack share show demo-share -f table
```

3) Allow access to a VM IP (adjust IP)

```
openstack share access allow demo-share ip 192.168.1.10
```

4) Mount in the VM

```
sudo apt-get update && sudo apt-get install -y nfs-common  
sudo mkdir -p /mnt/share  
# Replace <MANILA_EXPORT> with 'server:/export/path' from share show  
sudo mount -t nfs <MANILA_EXPORT> /mnt/share
```

Thuật ngữ: share = chia sẻ tệp; export (NFS) = điểm xuất khẩu; allow access = cấp quyền truy cập

Security & Operations Tips (Block/Object/File)

Block (Cinder)

- Prefer *encrypted volume types*; snapshot/backup routinely; detach before delete.
- Keep volume & VM in same AZ unless replication is configured.

Object (Swift)

- Enable versioning/lifecycle; use *temp URLs* for limited sharing; consider CDN.
- Set metadata: content-type, cache-control for static hosting.

File (Manila)

- Limit by IP/user/AD; per-share quotas; monitor throughput/latency.
- For many writers, validate file-locking vs. app semantics.

Thuật ngữ: encrypted volume = *đĩa mã hóa*; lifecycle policy = *chính sách vòng đời*; quota = *hạn ngạch*

Review OpenStack case study: Quick Comparison

Service	Model	Access/Protocol	Best For	Watch Out
Cinder	Block	iSCSI/NVMe (to host)	DBs, boot/system, low latency	AZ/attach affinity; snapshot backups
Swift	Object	HTTP REST (SDK/CLI)	Backups, logs, media, data lake	Higher latency; not POSIX; class consistency
Manila	File	NFS/SMB (multi-host)	Shared content, home dirs	Throughput tiering; file locking semantics

Thuật ngữ: iSCSI = kết nối khối qua IP; NVMe = giao diện lưu trữ nhanh; POSIX = chuẩn hệ tệp

Enhanced Tasks

- ① Cinder: create 10GiB volume; attach to VM; format & mount at /data.
- ② Swift: create container; upload a file; verify via CLI and show URL/endpoint.
- ③ Manila: create NFS share; allow VM IP; mount at /mnt/share; write/read test.

VM vs. Container Architecture

Aspect	Virtual Machine (VM)	Container
Kernel and Isolation	Each VM has its own guest kernel; isolation boundary is the <i>hypervisor</i> .	Share the host kernel; isolation via <i>namespaces</i> and <i>cgroups</i> , hardened by capabilities, seccomp, and <i>LSMs</i> .
Provisioning and Boot	High provisioning cost (allocate full OS, boot kernel); startup in seconds–minutes.	Lightweight provisioning; startup in milliseconds–seconds.
Filesystem (Overlay; CoW)	Independent disk images (e.g., QCOW2, VMDK).	Union/Copy-on-Write layers (e.g., overlay2); layer caching; use <i>bind mounts</i> or <i>block volumes</i> for write-heavy I/O.
Performance vs. Security	Strong isolation and stability, but higher resource overhead.	Higher density and efficiency; larger shared-kernel attack surface — requires defense-in-depth.

Cloud Application Isolation Models (I)

Workload Type	How it is Deployed	Isolation & Scalability	Key Drawbacks
Monolithic	Whole stack in one unit (single VM / host).	Simple deployment; vertical scaling.	Tightly coupled; difficult partial updates; poor fault isolation.
System-Containerized	System containers (LXC, systemd-nspawn) provide a full userspace.	Better process isolation than raw services; suitable for legacy stacks.	Heavier than app containers; limited orchestration support; shared kernel.
Application-Containerized (OCI)	Microservices packaged as OCI images; run via container runtimes and orchestrators.	Fast startup, horizontal scaling, CI/CD friendly, fine-grained isolation.	State must be externalized; kernel exposure risk; networking/storage integration complexity.

Thuật ngữ: OCI = chuẩn container mở; externalize = tách trạng thái ra ngoài

Cloud Application Isolation Models (II)

Model	Deployment / Runtime	Advantages	Limitations
Micro-container / Sandboxed	Micro-VMs or syscall sandboxes (Firecracker, gVisor, Kata).	VM-like isolation with near-container performance; good for multi-tenant workloads.	Increased complexity and some performance overhead.
Ephemeral / FaaS	Short-lived containers (event-driven).	Extreme elasticity; cost-efficient for bursty tasks.	Cold-start latency; unsuitable for long-running stateful jobs.

Progression: Installed Host → Isolated Service → Containerized App → Sandboxed Micro-container → Ephemeral Function

Thuật ngữ: ephemeral = tạm thời; sandbox = môi trường cách ly

Isolation Boundaries (Security & Fault Containment)

- **Process vs VM boundary:** containers = shared kernel (process-level). For untrusted tenants use micro-VMs / Kata / Firecracker.
- **Namespace isolation:** PID/NET/MNT/IPC/UTS/User separate process trees, network stacks, mounts, and identities.
- **Resource isolation:** enforce cgroups limits (cpu/memory/io/pids) to avoid noisy-neighbor and DoS.
- **Privilege reduction:** drop unneeded capabilities; run non-root user; use seccomp filters; enable AppArmor/SELinux; mount rootfs read-only.
- **Network isolation:** per-namespace veths, network policies, vSwitch security groups (OVS/OVN), and optional eBPF packet/flow filtering.
- **Storage isolation:** place secrets on tmpfs or secret stores; mount volumes with `ro`, `nodev`, `nosuid`, `noexec` where applicable.
- **Fault containment:** use liveness/readiness probes, restart policies, and non-shared PID namespaces to limit propagation of failures.

Thuật ngữ: seccomp = lọc syscall; capabilities = tập quyền Linux; eBPF = bộ lọc gói mở rộng

Scalability Mechanics (Density & Elasticity)

- **Horizontal scale (replicas):** immutable images allow fast cloning; stateless tiers scale linearly behind a load balancer.
- **Externalize state:** keep compute ephemeral; use managed DBs, object storage, and caches for session/state persistence.
- **Bin-packing & QoS:** declare resource requests/limits; use scheduler policies and QoS classes to meet p95/p99 latency SLOs.
- **Deployment strategies:** rolling, blue-green, and canary deployments to reduce blast radius during updates.
- **Autoscaling signals:** CPU/memory, queue length, RPS, latency percentiles; support scale-to-zero where appropriate.
- **Failure domains:** disperse replicas across hosts/availability zones; prefer anti-affinity for critical services.
- **Density vs safety:** increase density with quotas, limits and monitoring to avoid resource contention and noisy neighbors.

Thuật ngữ: SLO = mục tiêu mức dịch vụ; replica = bản sao

Execution Substrates

Substrate	How it Runs	When to Use	Notes
Linux host (containerd/runc)	Containers share host kernel; namespaces + cgroups + LSMs.	High-density stateless services; fast start.	Harden with non-root, seccomp, capability drops.
Rootless containers	User-namespace remapping, no host root.	Multi-tenant dev/CI and lower-privilege environments.	Feature limitations (net namespace, privileged operations).
gVisor / Firecracker	Syscall interception (gVisor) or micro-VMs (Firecracker).	Multi-tenant and untrusted workloads.	Stronger isolation; moderate performance overhead.
Kata / light VMs	Each pod runs in a minimal VM.	Zero-trust tenants, PCI/regulated workloads.	Lower density, longer cold-starts.
GPU / NUMA tuned nodes	Pin CPU/GPU and tune hugepages, I/O isolation.	ML training/inference, HPC.	Scheduler coordination needed (device plugins / topology hints).

Thuật ngữ: NUMA(Non-uniform memory access) = bộ nhớ bất đồng nhất; device plugin = trình cảm thiết bị cho scheduler

Container Stack, OCI and Supply Chain Security

- **Image model:** layered images, manifests, digests (sha256), tags, and multi-arch manifests.
- **Registry and policies:** private registries (Harbor, ECR, GHCR) with admission/gate policies, TTLs, and vulnerability scanning.
- **Runtime alignment:** prefer CRI-compatible runtimes (containerd / CRI-O) and validated low-level engines (runc, crun).
- **Supply-chain controls:** sign images (cosign), publish SBOMs (syft), and require attestation/provenance before deployment.
- **CI/CD integration:** automated build/test/sign → push → gated deploy pipeline with automated rollbacks on failures.

Thuật ngữ: SBOM(Software Bill of Materials) = danh sách thành phần phần mềm; cosign = ký số image

Practical Hardening — Image & Execution

- **Least privilege (runtime):** run as non-root; drop Linux capabilities; read-only rootfs; enable no-new-privileges.
- **Syscall surface:** attach a minimal seccomp profile per workload class.
- **LSM enforcement:** AppArmor/SELinux in enforcing mode; mask/read-only sensitive paths; restrict device mounts.
- **Image provenance & SBOM:** pin @sha256 digests; sign & verify images; publish SBOMs; scan before run.

Thuật ngữ: seccomp= lọc syscall; LSM (Linux Security Modules)= mô-đun bảo mật Linux; SBOM=danh sách thành phần phần mềm

Practical Hardening — Runtime & Operations

- **Runtime choice by risk:** containerd+runc for trusted/high-performance; gVisor/Kata for untrusted multi-tenant workloads.
- **Resource controls (cgroup v2):** cpu.max/weight, memory.high/max, pids.max, I/O throttling; (on K8s) set requests/limits & QoS to protect SLOs.
- **Network & secrets:** default-deny policy; constrain egress; eBPF filters for lateral-movement detection; use Vault or KMS-backed secrets; mount secrets as files (tmpfs).
- **Observability & resilience:** immutable logs; metrics & audit (seccomp/LSM denials); alert on anomalous syscalls/spikes; health probes, canary/rolling, auto-rollback.

Thuật ngữ: cgroup v2=nhóm điều khiển phiên bản 2; KMS (Key Management Service)=dịch vụ quản lý khóa; SLO (Service Level Objective)=mục tiêu mức dịch vụ

Trade-offs and Decision Guidance

- **Performance vs Isolation:** pure containers give best density/perf; micro-VMs/sandboxes give stronger isolation at cost of density.
- **Operational complexity:** stronger isolation increases surface area for ops (image signing, attestation, orchestration complexity).
- **When to choose what:**
 - High-trust internal services → containerd/runc on hardened hosts.
 - Multi-tenant or untrusted code → gVisor / Kata / Firecracker.
 - Event-driven short jobs → FaaS / ephemeral containers.
- **Rule of thumb:** start with secure defaults, progressively relax only when evidence (benchmarks / telemetry) justifies it.

Thuật ngữ: ephemeral = tạm thời; FaaS = Function-as-a-Service

runc

- **runc** is a lightweight, portable container runtime and the reference implementation of the OCI Runtime Spec.
- Originally extracted from Docker's *libcontainer*, now maintained under the **opencontainers** (OCI) org.
- Executes containers by setting up:
 - Linux **namespaces** (PID, network, mount, etc.)
 - **cgroups** for resource limits
 - Container filesystem and process entrypoint
- Invoked by higher-level runtimes (e.g., containerd, CRI-O) to spawn containers.

Thuật ngữ: runc = lớp runtime OCI thấp nhất, trực tiếp tạo tiến trình cách ly

containerd

- **containerd** is a **daemon-level runtime** that manages container lifecycle.
- Originated from Docker; now a **CNCF** project.
- Provides:
 - Image pull/push and **content store**
 - Snapshot management (e.g., overlayfs, btrfs, zfs)
 - Create/start/stop via OCI runtimes (e.g., runc, crun)
- No built-in networking; integrates with **CNI** via higher layers (e.g., Kubernetes CRI plugin).

Thuật ngữ: containerd = runtime trung gian điều phối ảnh, snapshot, và gọi runc/crun

CRI-O

- **CRI-O** is a lightweight runtime designed specifically for **Kubernetes**.
- Implements the **Container Runtime Interface (CRI)** natively.
- Uses:
 - runc or crun for low-level execution (OCI runtime).
 - common to supervise container processes and I/O.
- Image and storage via containers/image and containers/storage libraries (Podman/Buildah are sibling tools, not dependencies).
- Provides a minimal, secure, and fast runtime for Kubernetes nodes.

Thuật ngữ: CRI-O = runtime cho Kubernetes (CRI), dùng runc/crun + common

crun

- **crun** is a fast, lightweight alternative to **runc**.
- Written in C (vs Go for **runc**), often lower memory and faster startup.
- Fully compliant with the OCI Runtime Specification.
- Common in Fedora/RHEL; widely used with CRI-O.

Thuật ngữ: crun = thay thế runc, tối ưu hiệu năng/bộ nhớ

gVisor

- **gVisor** is a user-space kernel for containers (Google).
- Intercepts syscalls (via `runsc`) to sandbox workloads—adds an extra isolation layer.
- Trade-off: lower performance vs native `runc`.
- Suited for multi-tenant/untrusted workloads (e.g., GKE Sandbox, Cloud Run).

Thuật ngữ: gVisor = runtime an toàn, mô phỏng/syscall trong user space

Kata Containers

- Combines **lightweight VMs** with container interfaces (via a VMM such as KVM/Firecracker/QEMU).
- Each container runs inside a minimal VM → strong isolation boundary.
- Integrates with `containerd` and `CRI-O`.
- Used in multi-tenant cloud environments for stronger isolation.

Thuật ngữ: Kata = hybrid: container trong VM nhỏ, bảo mật cao

Docker Engine (Runtime Layer)

- **Docker Engine** includes:
 - dockerd — daemon/API
 - containerd — lifecycle management
 - runc — low-level OCI runtime
- Handles image build/pull (BuildKit) and local container lifecycle.
- Still widely used outside Kubernetes for standalone deployments.

Thuật ngữ: Docker Engine = dockerd + containerd + runc

Docker

- **Dockerfile:** FROM/RUN/COPY/ENV/WORKDIR/USER/EXPOSE/HEALTHCHECK/CMD.
- **CLI:** build/run/exec/ps/logs/push/pull; pin @sha256 instead of floating tags.
- **Best practices:** multi-stage builds; non-root USER; distroless; cache layers; HEALTHCHECK.
- **Secrets/Config:** never bake secrets; prefer files mounted from a secret store or Docker/K8s secrets; avoid plain env for sensitive data.

Thuật ngữ: distroless = ảnh tối giản không shell; pin digest = ghim theo băm

Container Platform — Overview

- A **container platform** provides full lifecycle for *single-host* containers:
 - Build, store, run, and secure containers (optionally compose multi-container apps).
 - Integrates with *orchestrators* for cluster scheduling/scaling; ties into CI/CD and monitoring.
- Built **on top of container runtimes** (e.g., containerd, CRI-O).
- Examples: Docker, Podman, OpenShift (platform product built on Kubernetes).

Thuật ngữ: Platform = tầng quản lý build/run/single-host, dựa trên runtime; scale cụm do orchestrator

Platform vs Orchestration (Concept)

- **Container Platform:** developer/runtime tooling for *build, image management, local run, developer workflows* (CLI, build, registry, compose). Examples: Docker, Podman, OpenShift (platform product).
- **Orchestration:** cluster control plane that *schedules, scales, networks, discovers, and heals* workloads across many hosts. Examples: Kubernetes, Docker Swarm (legacy).
- **Short rule:** Platform = “build + ship + run (single host/node)” — Orchestration = “schedule + scale + operate (cluster)”.

Thuật ngữ: Platform = nền tảng; Orchestration = điều phối cụm

Docker — Orchestration (note)

- **Docker Swarm:** built-in orchestrator mode in Docker Engine (services, simple scheduling) — now legacy/maintenance.
- **Positioning:** Swarm is an orchestration feature, not the same as Docker the platform. Modern clusters predominantly use Kubernetes.
- **Implication:** when you say "Docker", be explicit: developer platform (Engine/Compose) vs. orchestration mode (Swarm).

Thuật ngữ: Swarm = chế độ điều phối của Docker (legacy); phân biệt rõ để tránh nhầm lẫn

Comparison — Platform vs Orchestrator

Role	Platform (build/run/dev)	Orchestrator (cluster operate)
Primary concerns	developer workflows, image build, local run, registries, single-node lifecycle	scheduling, scaling, service discovery, cluster networking, reconciliation
Examples	Docker Engine, Docker Desktop, Podman, OpenShift (platform product)	Kubernetes (kube-apiserver, scheduler), Docker Swarm (legacy)
Typical outputs	images, Compose files, OCI artifacts	Pods, Deployments, Services, NetworkPolicies
Interaction	runs workloads on a node via a runtime (containerd/runc)	delegates execution to CRI runtimes across nodes

Thuật ngữ: Nói rõ “platform” hay “orchestrator” khi thiết kế kiến trúc

Example: Multi-stage Dockerfile (Node.js)

```
# Build stage (needs dev deps)
FROM node:20-alpine AS build
WORKDIR /app
COPY package*.json ./
RUN npm ci
COPY . .
RUN npm run build
# Runtime stage (production deps only)
FROM gcr.io/distroless/nodejs20-debian12
WORKDIR /app
COPY --from=build /app/dist ./dist
# Option B (simpler): add an intermediate "deps" stage to produce prod-only node_modules.
USER 10001:10001
ENV NODE_ENV=production
EXPOSE 3000
CMD ["dist/server.js"]
```

Thuật ngữ: multi-stage build = xây nhiều giai đoạn; runtime image = ảnh chạy

Container Networking Basics

- **Drivers:** bridge, host, macvlan, overlay (CNI concepts).
- **Port mapping:** -p host:container; DNAT via iptables/nft; outbound SNAT.
- **Service discovery:** embedded DNS; Compose/K8s Services; health/readiness probes.
- **North/South vs. East/West:** publish/LB vs. service-to-service mesh/policies.

Thuật ngữ: CNI = giao diện mạng container; DNAT/SNAT = dịch đích/nguồn

Container Storage & Configuration

- **Volumes:** bind mounts vs. named volumes; local/NFS/CIFS; avoid heavy writes on overlay2.
- **State:** separate data to persistent volumes; use snapshot/backup cycles.
- **Secrets/Config:** Swarm/K8s Secrets; never commit secrets; rotate periodically.

Thuật ngữ: bind mount = *gắn kết đường dẫn*; persistent volume = *độ bền vững*; rotate = *xoay vòng*

Platform Comparison

Runtime/Tool	Strengths	Notes
Docker Engine	Rich ecosystem, UX, Compose	K8s prod uses containerd/CRI; rootless available
containerd	CRI standard, light, K8s-native	Needs CLI/UX (nerdctl)
CRI-O	K8s-focused, minimal surface	Narrow scope (K8s-centric)
Podman	Daemonless, rootless, Docker-like CLI	Compose via podman-compose

Thuật ngữ: rootless = không đặc quyền root; daemonless = không nền dịch vụ

Operational Considerations & Security

- **Provenance:** cosign signatures, SBOM (syft), attestations; reject unsigned images.
- **Least privilege:** non-root USER; drop caps; seccomp; AppArmor/SELinux; read-only FS.
- **Resource limits:** -cpus, -memory, -pids-limit; avoid noisy neighbors/DoS.
- **Policies:** admission control (OPA/Conftest); pin-by-digest; base-image allowlists.

Thuật ngữ: attestation = xác nhận tạo tác; allowlist = danh sách cho phép; admission control = kiểm soát nhập cụm

Practice Overview

- **Task 1:** Provision OpenStack VM; install Docker via Ansible (idempotent).
- **Task 2:** Build multi-stage image; push to private registry.
- **Task 3:** Run container with limits; health checks; non-root user.
- **Task 4 (optional):** Integrate Neutron SGs, Cinder volumes; optionally Magnum/Zun.

Thuật ngữ: idempotent = *bất biến theo lắp*; registry = *kho ảnh*

Ansible Snippet: Install Docker on OpenStack VM

```
- hosts: app_vms
become: true
tasks:
  - name: Install prerequisites
    package:
      name: [curl, ca-certificates, gnupg, lsb-release]
      state: present

  - name: Install Docker Engine
    shell: |
      install -m 0755 -d /etc/apt/keyrings
      curl -fsSL https://download.docker.com/linux/$(. /etc/os-release; echo $ID)/gpg | \
        gpg --dearmor -o /etc/apt/keyrings/docker.gpg
      echo \
        "deb [arch=$(dpkg --print-architecture) signed-by=/etc/apt/keyrings/docker.gpg] \
        https://download.docker.com/linux/$(. /etc/os-release; echo $ID) \
        $(. /etc/os-release; echo $VERSION_CODENAME) stable" | \
      tee /etc/apt/sources.list.d/docker.list > /dev/null
    apt-get update
    apt-get install -y docker-ce docker-ce-cli containerd.io
```

Ansible Snippet: Build & Push Image

```
- hosts: app_vms
become: true
vars:
  registry: "ghcr.io/yourorg"
  image_name: "wk7-sample"
  image_tag: "v1"
tasks:
  - name: Copy app sources
    synchronize:
      src: ./app/
      dest: /opt/app/

  - name: Build image
    community.docker.docker_image:
      name: "{{ registry }}/{{ image_name }}:{{ image_tag }}"
      build:
        path: /opt/app
      push: no

  - name: Login and push
    community.docker.docker_login:
```

Ansible Snippet: Run Container with Limits

```
- hosts: app_vms
become: true
tasks:
  - name: Run container with resource limits
    community.docker.docker_container:
      name: wk7-web
      image: "ghcr.io/yourorg/wk7-sample:v1"
      state: started
      published_ports:
        - "80:3000"
      restart_policy: unless-stopped
      memory: "512m"
      cpus: "0.5"
      env:
        NODE_ENV: production
```

Thuật ngữ: published ports = cổng công bố; restart policy = chính sách khởi động lại

Integrating with OpenStack

- **Neutron:** SG mở 80/443; Floating IP để publish; inbound via LBaaS khi cần.
- **Cinder:** mount volume làm thư mục dữ liệu (-v /data:/var/lib/app); snapshot/backup.
- **Glance/Packer:** base image đã cài Docker để boot nhanh và nhất quán.
- **Magnum/Zun (tùy chọn):** Container/K8s as-a-service trong OpenStack.

Thuật ngữ: SG = nhóm bảo mật; LBaaS = cân bằng tải như dịch vụ

Mini Lab (Tóm tắt)

- ① Cài Docker trên VM OpenStack; cấu hình SG cho 22/80.
- ② Build image multi-stage; push lên registry riêng (pin digest).
- ③ Chạy container non-root; giới hạn CPU/RAM; thiết lập HEALTHCHECK.
- ④ Cấp Floating IP; kiểm tra HTTP từ Internet.
- ⑤ (Tuỳ chọn) Mount Cinder volume cho dữ liệu; thử backup/restore nhanh.

Thuật ngữ: backup/restore = sao lưu/khôi phục; non-root = không đặc quyền root

Assessment & Deliverables

- **Báo cáo (2 trang)**: so sánh VM vs. container (perf, security), kiến trúc triển khai, kết quả kiểm thử.
- **Minh chứng**: ảnh cài Docker; ảnh image trong registry; docker ps; endpoint công khai.
- **File nộp**: Dockerfile, playbooks Ansible, inventory, compose (nếu dùng).
- **Rubric (10 điểm)**: kỹ thuật (4), bảo mật (3), tự động hóa (2), tài liệu (1).

Thuật ngữ: inventory = danh mục máy; endpoint = điểm cuối

Threat Model Lite (STRIDE) cho Containers

- **Spoofing:** rò rỉ registry creds ⇒ secret-manager/KMS, token scope hép.
- **Tampering:** sửa image ⇒ ký số, pin digest, admission policy.
- **Repudiation:** thiếu audit ⇒ log build/push/pull, provenance/SBOM.
- **Info Disclosure:** secrets trong image ⇒ build-args/secret mounts.
- **DoS:** hog CPU/RAM ⇒ -cpus, -memory, -pids-limit.
- **EoP:** -privileged ⇒ tránh; drop caps; seccomp profile chặt.

Thuật ngữ: KMS = hệ quản lý khoá; EoP = leo thang đặc quyền

Pitfalls & Anti-patterns

- **Floating tags:** dùng :latest gây drift ⇒ luôn pin @sha256.
- **Run as root:** root + volume RW ⇒ tăng rủi ro EoP.
- **overlay2 write-heavy:** workload I/O ngẫu nhiên ⇒ chuyển sang volumes chuyên dụng.
- **Bloat images:** base image quá lớn ⇒ multi-stage + distroless.

Thuật ngữ: drift = trôi lệch cấu hình; RW = đọc-ghi

Wrap-up

- Cloud host workloads differ by **isolation boundary** and **scalability mechanics**.
- SDN/Storage choices shape connectivity & persistence; container practice shapes density & speed.
- Safe defaults: non-root, signed images, SBOM, pin digest, strict limits, secrets outside images.
- Next: **Orchestration with Kubernetes (K8s basics)**.

Thuật ngữ: isolation boundary = ranh giới cô lập; density = mật độ